

CSC718 Parallel Programming Midterm

Total: 100 points

Due by Midnight, Oct 16, 2022

(Midterm is a take-home exam. It includes four questions. The last one is a programming assignment. Please submit all your solutions and source code through D2L midterm Dropbox.)

- 1) (20 points) Binary semaphores only allow the integer to hold the values 0 and 1. However, binary semaphore are easier to implement than general semaphores. Here we demonstrate that a general semaphore can be implemented from a binary semaphore.
- A structure of general semaphore is defined as below. It utilizes two binary semaphores. The binary semaphores support two functions.
- b_wait(s): decrements s if s>0. If not, the process executing is blocked.
- b_signal(s): increments s by 1 (does not change the value if s = 1). After the increment, a process waiting on s resumes its execution.

```
struct semaphore
{
    int value;
    int waiting = 0;
    b_semaphore queue = 0; // initialized to 0
    b_semaphore mutex = 1; // initialized to 1
}

void wait(struct semaphore *s)
{
    b_wait(s->mutex);
    if (s->value == 0)
    {
        s->waiting++;
        b_signal(s->mutex);
        b_wait(s->queue);
    }
    else
    {
        s->value--;
        A. _____
    }
}

void signal(struct semaphore *s)
{
    b_wait(s->mutex);
    if (s->waiting > 0)
    {
        s->waiting--;
        B. _____
        C. _____
    }
    else
    {
        s->value++;
        b_signal(s->mutex);
    }
}
```

Figure 2

Two semaphore functions, wait(s) and signal(s) are given in Figure 2. However, there are three lines of code are missing. Fill the missing code in A, B, and C.

A. b_signal(s->mutex);

- B. b_signal(s->mutex);
 C. b_signal(s->queue);

2) (24 points) Assume n pieces of work forming into a one-dimensional array (0 to $n-1$), p processes, and block allocation (scatters larger blocks among processes) is used. Process rank i starts from 0.

- a. What is the most pieces of work any process has?

$$\left\lceil \frac{n}{p} \right\rceil$$

- b. What is the least pieces of work any process has?

$$\left\lfloor \frac{n}{p} \right\rfloor$$

- c. How many processes have the most pieces of work?

$$r = n \bmod p$$

- d. What is the first element controlled by process i ?

$$\left\lfloor \frac{in}{p} \right\rfloor$$

- e. What is the last element controlled by process i ?

$$\left\lfloor \frac{(i+1)n}{p} \right\rfloor - 1$$

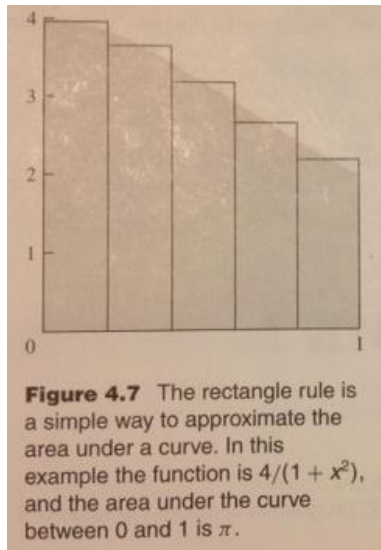
- f. Which process will have the control on element j ?

$$\left\lfloor \frac{(p(j+1) - 1)}{n} \right\rfloor$$

3) (36 points) Programming Assignment: The value of the definite integral

$$\int_0^1 \frac{4}{1+x^2} dx$$

is π . We can use numerical integration to compute π by approximating the area under the curve. A simple way to do this is called the rectangle rule (Figure 4.7). We divide the interval $[0, 1]$ into k subintervals of equal size. We find the height of the curve at the midpoint of each of these subintervals. With these heights we can construct k rectangles. The area of the rectangles approximates the area under the curve. As k increases, the accuracy of the estimate also increases.



A C program that uses the rectangle rule to approximate π appears in Figure 4.8. The source code, pi.c, can be found in the class website. You can compile and run the program in Linux.

```
#include <stdio.h>
#define INTERVALS 1000000
int main(int argc, char* argv[])
{
    double area; /* The final answer */
    double ysum; /* Sum of rectangle heights */
    double xi; /* Midpoint of interval */
    int i;
    ysum= 0.0;
    for(i=0;i<INTERVALS;i++)
    {
        xi= ((1.0/INTERVALS)*(i+0.5)); /* Midpoint of interval */
        ysum += 4.0/(1.0 + xi*xi);
    }
    area=ysum*(1.0/INTERVALS);
    printf("Area is %13.11fn",area);
    return 0;
}
```

Figure 4.8 A C program to compute the value of π using the rectangle rule.

- a. (10 points) Let $\text{INTERVALS}=1000000$, what is pi? Show your calculation result.

$$\int_0^1 \frac{4}{1+x^2} dx = 4 \int_0^1 \frac{1}{1+x^2} dx = 4 \tan^{-1} x + C$$

- pi is 3.14159265359

```
1      #include <stdio.h>
2      #define INTERVALS 1000000
3
4      int main(int argc, char* argv[])
5      {
6          double area;    /* The final answer */
7          double ysum;    /* Sum of rectangle heights */
8          double xi;      /* Midpoint of interval */
9          int i;
10
11         ysum = 0.0;
12
13         for (i=0; i < INTERVALS; i++)
14         {
15             xi=((1.0/INTERVALS)*(i+0.5));
16             ysum+=4.0/(1.0+xi*xi);
17         }
18
19         area = ysum * (1.0/INTERVALS);
20         printf("pi is %13.11f\n", area);
21         return 0;
22     }
```

main

Run: scratch x

C:\cPrograms\CLionProjects\scratch\cmake-build-debug\scratch.exe
pi is 3.14159265359

Process finished with exit code 0

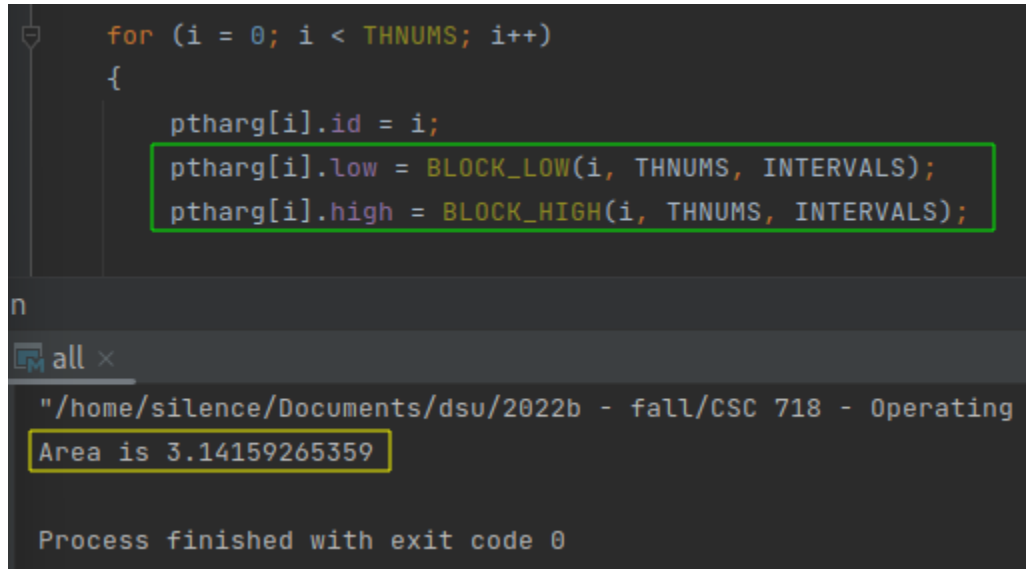
- b. (16 points) The program can be converted to a multithread program. A sample program, mth-pi.c, has been provided as below. There are two lines, A and B, which are incomplete. Fill the blanks and complete the program.

ptharg[i].low = _____ A _____;
ptharg[i].high = _____ B _____;

- A. BLOCK_LOW(i, THNUMS, INTERVALS)
B. BLOCK_HIGH(i, THNUMS, INTERVALS)

- c. (10 points) Fill the missing lines in mth-pi.c. Compile and run the program. Let INTERVAL=1000000, what is pi? Is it the same value as your calculation result in a).

- pi is 3.14159265359, it is the same value as my calculation result in a



The screenshot shows a code editor with a C program. The code defines a struct `ThreadParams` with fields `id`, `low`, `high`, and `ysum`. It also defines macros for `BLOCK_LOW`, `BLOCK_HIGH`, and `BLOCK_SIZE`. The `main` function sets `THNUMS` to 3 and `INTERVALS` to 1000000. It creates an array of `ThreadParams` and fills it with values. The `pthread_t` array is then used to create threads. The terminal output shows the program's execution, including the path to the source file and the calculated value of pi: "Area is 3.14159265359". The process finished with exit code 0.

```
for (i = 0; i < THNUMS; i++)
{
    ptharg[i].id = i;
    ptharg[i].low = BLOCK_LOW(i, THNUMS, INTERVALS);
    ptharg[i].high = BLOCK_HIGH(i, THNUMS, INTERVALS);
}

n

all x
"/home/silence/Documents/dsu/2022b - fall/CSC 718 - Operating
Area is 3.14159265359

Process finished with exit code 0
```

```
#include <pthread.h>
#include <stdio.h>
#define INTERVALS 1000000
```

```
#define THNUMS      3
```

```
// thread parametes
struct ThreadParams
{
    int id;           // id
    int low;          // start
    int high;         // end
    double ysum;      // return partial sum
};
```

```
#define BLOCK_LOW(id, p, n) ((id)*(n)/(p))
#define BLOCK_HIGH(id, p, n) (BLOCK_LOW((id)+1,p,n)-1)
#define BLOCK_SIZE(id, p, n) (BLOCK_HIGH(id, p, n)-BLOCK_LOW(id, p, n)+1)
```

```
// calculate pi partial sum
void *calcp_i(void *arg)
{
    double ysum;
    double xi;
    int i;
```

```

    struct ThreadParams *pm = (struct ThreadParams*)arg;

    ysum = 0.0;

    for (i = pm->low; i <= pm->high; i++)
    {
        xi=((1.0/INTERVALS)*(i+0.5));
        ysum=ysum+4.0/(1.0+xi*xi);
    }

    pm->ysum = ysum;
}

int main(int arc, char* argv[])
{
    struct ThreadParams ptharg[THNUMS];
    double area;
    double ysum;
    int i;
    pthread_t tid;
    void *status;

    // create multithreads to calculate partial sum
    ysum = 0.0;

    for (i = 0; i < THNUMS; i++)
    {
        ptharg[i].id = i;
        ptharg[i].low = _____ A _____;
        ptharg[i].high = _____ B _____;

        if (pthread_create(&tid, NULL, calcpi, (void*)&ptharg[i]) != 0)
        {
            perror("error creating child");
            return -1;
        }

        pthread_join(tid, &status);
    }

    // calculate total area
    area = 0;

    for (i = 0; i < THNUMS; i++)

```

```

{
    area = area + ptharg[i].ysum;
}

area = area * (1.0/INTERVALS);
printf("Area is %13.11f\n", area);
return 0;
}

```

- 4) (20 points) Programming Assignment: Write a parallel program using MPI that computes the sum $1+2+\dots+p$ in the following manner: Each process i assigns the value $i+1$ to an integer, and then the processes perform a sum reduction of these values. Process 0 should print the result of the reduction. As a way of double checking the result, process 0 should also compute and print the value $p(p+1)/2$. Please benchmark the running time of the parallel program using the Rushmore cluster as discussed in the class.

	Np=1	Np=2	Np=3	Np=4
Execution time	0.000020	0.000165	0.000111	0.000281

```

mpiuser@Jefferson:~/Desktop/KC$ mpicc sum.c -o sum
mpiuser@Jefferson:~/Desktop/KC$ scp /home/mpiuser/Desktop/KC/sum mpiuser@Washington:/home/mpiuser/Desktop/KC/sum
sum                               100% 17KB 13.2MB/s 00:00
mpiuser@Jefferson:~/Desktop/KC$ scp /home/mpiuser/Desktop/KC/sum mpiuser@Lincoln:/home/mpiuser/Desktop/KC/sum
sum                               100% 17KB 14.5MB/s 00:00
mpiuser@Jefferson:~/Desktop/KC$ scp /home/mpiuser/Desktop/KC/sum mpiuser@Roosevelt:/home/mpiuser/Desktop/KC/sum
sum                               100% 17KB 13.6MB/s 00:00
mpiuser@Jefferson:~/Desktop/KC$ mpirun -np 1 -machinefile machinefile.dsu ./sum
Lincoln [Processor 0 of 1] has completed.

The sum of all values between 1 and 10 is 55
Elapsed time: 0.000020.
mpiuser@Jefferson:~/Desktop/KC$ mpirun -np 2 -machinefile machinefile.dsu ./sum
Washington [Processor 1 of 2] has completed.
Lincoln [Processor 0 of 2] has completed.

The sum of all values between 1 and 10 is 55
Elapsed time: 0.000165.
mpiuser@Jefferson:~/Desktop/KC$ mpirun -np 3 -machinefile machinefile.dsu ./sum
Washington [Processor 1 of 3] has completed.
Roosevelt [Processor 2 of 3] has completed.
Lincoln [Processor 0 of 3] has completed.

The sum of all values between 1 and 10 is 55
Elapsed time: 0.000111.
mpiuser@Jefferson:~/Desktop/KC$ mpirun -np 4 -machinefile machinefile.dsu ./sum
Jefferson [Processor 3 of 4] has completed.
Washington [Processor 1 of 4] has completed.
Roosevelt [Processor 2 of 4] has completed.
Lincoln [Processor 0 of 4] has completed.

The sum of all values between 1 and 10 is 55
Elapsed time: 0.000281.

```