

## CSC718 Parallel Programming Midterm

Total: 100 points

Due by Midnight, Oct 16, 2022

(Midterm is a take-home exam. It includes four questions. The last one is a programming assignment. Please submit all your solutions and source code through D2L midterm Dropbox.)

- 1) (20 points) Binary semaphores only allow the integer to hold the values 0 and 1. However, binary semaphore are easier to implement than general semaphores. Here we demonstrate that a general semaphore can be implemented from a binary semaphore.
- A structure of general semaphore is defined as below. It utilizes two binary semaphores. The binary semaphores support two functions.
- b\_wait(s): decrements s if s>0. If not, the process executing is blocked.
- b\_signal(s): increments s by 1 (does not change the value if s = 1). After the increment, a process waiting on s resumes its execution.

```
struct semaphore
{
    int value;
    int waiting = 0;
    b_semaphore queue = 0; // initialized to 0
    b_semaphore mutex = 1; // initialized to 1
}

void wait(struct semaphore *s)
{
    b_wait(s->mutex);
    if (s->value == 0)
    {
        s->waiting++;
        b_signal(s->mutex);
        b_wait(s->queue);
    }
    else
    {
        s->value--;
        A. _____
    }
}

void signal(struct semaphore *s)
{
    b_wait(s->mutex);
    if (s->waiting > 0)
    {
        s->waiting--;
        B. _____
        C. _____
    }
    else
    {
        s->value++;
        b_signal(s->mutex);
    }
}
```

Figure 2

Two semaphore functions, wait(s) and signal(s) are given in Figure 2. However, there are three lines of code are missing. Fill the missing code in A, B, and C.

A. \_\_\_\_\_

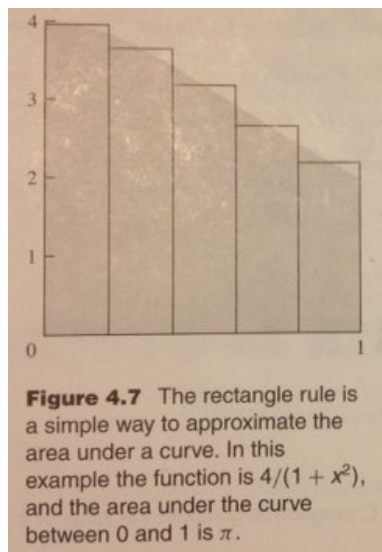
- B. \_\_\_\_\_  
C. \_\_\_\_\_

- 2) (24 points) Assume  $n$  pieces of work forming into a one-dimensional array (0 to  $n-1$ ),  $p$  processes, and block allocation (scatters larger blocks among processes) is used. Process rank  $i$  starts from 0.
- What is the most pieces of work any process has?
  - What is the least pieces of work any process has?
  - How many processes have the most pieces of work?
  - What is the first element controlled by process  $i$ ?
  - What is the last element controlled by process  $i$ ?
  - Which process will have the control on element  $j$ ?

- 3) (36 points) Programming Assignment: The value of the definite integral

$$\int_0^1 \frac{4}{1+x^2} dx$$

is  $\pi$ . We can use numerical integration to compute  $\pi$  by approximating the area under the curve. A simple way to do this is called the rectangle rule (Figure 4.7). We divide the interval  $[0, 1]$  into  $k$  subintervals of equal size. We find the height of the curve at the midpoint of each of these subintervals. With these heights we can construct  $k$  rectangles. The area of the rectangles approximates the area under the curve. As  $k$  increases, the accuracy of the estimate also increases.



A C program that uses the rectangle rule to approximate  $\pi$  appears in Figure 4.8. The source code, `pi.c`, can be found in the class website. You can compile and run the program in Linux.

```

#include <stdio.h>
#define INTERVALS 1000000
int main(int argc, char* argv[])
{
    double area; /* The final answer */
    double ysum; /* Sum of rectangle heights */
    double xi; /* Midpoint of interval */
    int i;
    ysum= 0.0;
    for(i=0;i<INTERVALS;i++)
    {
        xi= ((1.0/INTERVALS)*(i+0.5)); /* Midpoint of interval */
        ysum += 4.0/(1.0 + xi*xi);
    }
    area=ysum*(1.0/INTERVALS);
    printf("Area is %13.11f\n",area);
    return 0;
}

```

Figure 4.8 A C program to compute the value of  $\pi$  using the rectangle rule.

- (10 points) Let INTERVALS=1000000, what is pi? Show your calculation result.
- (16 points) The program can be converted to a multithread program. A sample program, mth-pi.c, has been provided as below. There are two lines, A and B, which are incomplete. Fill the blanks and complete the program.

```

ptharg[i].low = _____ A _____;
ptharg[i].high = _____ B _____;

```

- (10 points) Fill the missing lines in mth-pi.c. Compile and run the program. Let INTERVAL=1000000, what is pi? Is it the same value as your calculation result in a).

```

#include <pthread.h>
#include <stdio.h>
#define INTERVALS 1000000

```

```

#define THNUMS      3

```

```

// thread parametes
struct ThreadParams
{

```

```

    int id;                // id
    int low;               // start

```

```

        int high;                // end
        double ysum;             // return partial sum
};

#define BLOCK_LOW(id, p, n)    ((id)*(n)/(p))
#define BLOCK_HIGH(id, p, n)  (BLOCK_LOW((id)+1,p,n)-1)
#define BLOCK_SIZE(id, p, n)  (BLOCK_HIGH(id, p, n)-BLOCK_LOW(id, p, n)+1)

// calculate pi partial sum
void *calcp(void *arg)
{
    double ysum;
    double xi;
    int i;

    struct ThreadParams *pm = (struct ThreadParams*)arg;

    ysum = 0.0;

    for (i = pm->low; i <= pm->high; i++)
    {
        xi=((1.0/INTERVALS)*(i+0.5));
        ysum=ysum+4.0/(1.0+xi*xi);
    }

    pm->ysum = ysum;
}

int main(int arc, char* argv[])
{
    struct ThreadParams ptharg[THNUMS];
    double area;
    double ysum;
    int i;
    pthread_t tid;
    void *status;

    // create multithreads to calculate partial sum
    ysum = 0.0;

    for (i = 0; i < THNUMS; i++)
    {
        ptharg[i].id = i;
        ptharg[i].low = _____ A _____;
    }
}

```

```

    ptharg[i].high = _____B_____ ;

    if (pthread_create(&tid, NULL, calcpi, (void*)&ptharg[i]) != 0)
    {
        perror("error creating child");
        return -1;
    }

    pthread_join(tid, &status);
}

// calculate total area
area = 0;

for (i = 0; i < THNUMS; i++)
{
    area = area + ptharg[i].ysum;
}

area = area * (1.0/INTERVALS);
printf("Area is %13.11f\n", area);
return 0;
}

```

- 4) (20 points) Programming Assignment: Write a parallel program using MPI that computes the sum  $1+2+\dots+p$  in the following manner: Each process  $i$  assigns the value  $i+1$  to an integer, and then the processes perform a sum reduction of these values. Process 0 should print the result of the reduction. As a way of double checking the result, process 0 should also compute and print the value  $p(p+1)/2$ . Please benchmark the running time of the parallel program using the Rushmore cluster as discussed in the class.

|                | Np=1 | Np=2 | Np=3 | Np=4 |
|----------------|------|------|------|------|
| Execution time |      |      |      |      |